



Big Data BBQ

Dining at the OpenStack buffet to create cloud applications

Michael McCune
Senior Software Engineer
Red Hat

Nikita Konovalov
Software Engineer
Mirantis

April 28, 2016

Caveats

- This is not a blessing of one technology over another
- Some of the samples and features discussed may involve master branch or experimental source code, care should be taken when replicating it in production environments

Problem


- Want to create big data applications on OpenStack
- Want them to use cloud resources
- Don't want to bug operations team for continuous deployments

Solution


- Build cloud applications
- Create a resilient infrastructure
- Empower developers to run their own upgrades

Cloud applications


- Collections of applications
- Mixed types; console, server, graphical
- Resource controllers



Cloud applications




Cloud applications




Building a big data application

- What data will you process
- How will you process the data
- Where will you store the results



Planning

- Moving pieces
- The data pipeline
- Storage
- Custom applications



Resiliency

- Build composable pieces
- Be aware of network functionality
- Stateless applications



Martin Kalfatovi, CC BY-NC-SA 2.0

Empowering developers

- Appropriate resource levels
- Access to cloud services
- Production versus Development



Tools of the trade

- Version control
- Configuration management
- Test gating
- Python REPL

Sounds great, now what?

- Foundational elements
- Data processing clusters
- Message queues
- Data stores
- Spark components
- Custom pieces



Alan Chia, CC BY-SA 2.0

OpenStack tooling

- OpenStack Python clients
- Oslo projects
 - oslo.config
 - oslo.log
- Development environments
 - Devstack
 - Kolla
 - OpenStack-Ansible

Deployment

- Create custom images
- Modify running instances
- Re-deploy or update

Deployment

```
if [ "${DIB_DEBUG_TRACE:-0}" -gt 0 ]; then
    set -x
fi
set -eu
set -o pipefail

pushd /opt

echo "Cloning big-data-bbq"
git clone https://github.com/elmiko/big-data-bbq
echo "Installing requirements"
pip install -r big-data-bbq/requirements.txt

popd
```


Deployment

```
env.host_string='10.0.0.3'  
env.key_filename='/home/mike/.ssh/cloud_key'  
env.user='fedora'  
  
with hide('running', 'stdout', 'stderr'):  
    with cd('/opt'):  
        run('git clone http://github.com/elmiko/big-data-bbq')  
        sudo('pip install -r big-data-bbq/requirements.txt')
```

Configurations

- Command line parameters
- File based options
- Protecting sensitive information

Configurations

```
sahara_cluster_option = [  
    cfg.StrOpt('float_pool_network_name', default='public',  
               help='The name of the floating pool network to use. '),  
    cfg.StrOpt('management_network_name', default='private',  
               help='The name of the management network to use. '),  
]  
  
conf.register_opts(sahara_cluster_options)  
conf.register_opts(password.Password.get_options(), group='keystone')
```

Configurations

```
[DEFAULT]
```

```
# The name of the floating pool network to use. (string value)
```

```
#float_pool_network_name = public
```

```
# The name of the management network to use. (string value)
```

```
#management_network_name = private
```

```
[keystone]
```

```
# Authentication URL (string value)
```

```
auth_url = http://mystack.myhost.org:5000/v3
```

```
# User ID (string value)
```

```
user_id = elmiko
```

Authentication

- Common to all OpenStack components*
- Widely repeated code
- Reusable sessions

Authentication

```
auth = password.Password.load_from_conf_options(conf, group='keystone')
sesh = session.Session(auth=auth)

keystone_client = keystone.Client(session=sesh)

sahara_client = sahara.Client(session=sesh)

trove_client = trove.Client(conf.keystone.username,
                             conf.keystone.password,
                             conf.keystone.project_name,
                             conf.keystone.auth_url)
```

Data processing clusters

- Operation templates
- Provisioning and scaling
- Monitoring cluster progress

Data processing clusters

```
kwargs = {'name': name,  
          'plugin_name': 'spark',  
          'hadoop_version': '1.6.0',  
          'net_id': mgmt_net,  
          'cluster_configs': {'general': {'Enable NTP service': False}},  
          'node_groups': [  
            {'count': 1,  
             'name': 'spark160-master',  
             'flavor_id': '2',  
             'node_processes': ['namenode', 'master'],  
             'floating_ip_pool': float_pool  
            },  
            {'count': 3,  
             'name': 'spark160-worker',
```


Data processing clusters

```
template_id = sahara_client.cluster_templates.create(**kwargs).id

image_id = sahara_client.images.find(name='my_spark_image')[0].id

keypair_id = nova_client.keypairs.find(name='cloud_key').id

cluster = sahara_client.clusters.create(name='my spark cluster',
                                         plugin_name='spark',
                                         hadoop_version='1.6.0',
                                         cluster_template_id=template_id,
                                         default_image_id=image_id,
                                         user_keypair_id=keypair_id)
```

Data processing clusters

```
while cluster.status != 'Active':  
    time.sleep(5)  
    cluster = sahara_client.clusters.get(cluster.id)  
    if cluster.status == "Error":  
        raise Exception('cluster blew up!')  
  
master_ip = None  
for group in cluster.node_groups:  
    if group.get('name', '') == 'spark160-master':  
        master_ip = group.get('instances', [{}])[0].get('management_ip')
```

Message queues

- Named queues
- Writing messages, options
- Reading, affecting the queue
- Subscribing, how to receive

Message queues

```
# setup queue
queue = zaqar_client.queue('my-data-channel')

# send message
message = {'data': {'value': 9000, 'type': 'strength'}}
body = json.dumps(message)
payload = {'body': body, 'ttl': 60}
queue.send(payload)
```

Message queues

```
# setup queue
queue = zaqar_client.queue('my-data-channel')

# receive message
while True:
    messages = [m for m in queue.pop(count=10)]
    if len(messages) == 0:
        time.sleep(1)
        continue
    for msg in messages:
        process_message(msg.body)
```

Data stores

- What type of store
- Maintenance and longevity
- Providing access to data

Data stores

```
databases = [{'name': database_name}]
users = [{'name': username,
          'password': password,
          'databases': [{'name': database_name}]}]
nics = [{'net-id': network_id}]

instance = trove_client.instances.create(
    name=name, flavor_id=flavor, datastore=datastore_id,
    datastore_version=datastore_version_id, databases=databases,
    users=users, volume={'size': 10}, nics=nics)
```

Data stores

```
while instance.status != 'ACTIVE':
    time.sleep(5)
    instance = trove_client.instances.find(id=instance.id)
    if instance.status == 'ERROR':
        raise Exception('Oops, database instance create FAIL!')

ips = []
for ip in instance.ip:
    m = re.match('^[0-9]{1,3}\.[0-9]{1,3}\.[0-9]{1,3}\.[0-9]{1,3}$', ip)
    if m is not None:
        ips.append(m.group())

conxstr = 'mongodb://{user}:{passwd}@{ip}:27017/{database}'.format(
    username, password, ips[0], database_name)
```


Spark streaming

- Cluster connectivity
- Time slices
- Queues or direct communications

Spark streaming

```
args = parser.parse_args()

sconf = SparkConf().setAppName(args.appname)
if args.master:
    sconf.setMaster(args.master)
sc = SparkContext(conf=sconf)
ssc = StreamingContext(sc, 1)

lines = ssc.socketTextStream(args.socket, args.port)
lines.foreachRDD(lambda rdd: process_generic(rdd, args.mongo_url, args.rest_url))

ssc.start()
ssc.awaitTermination()
```

Spark streaming

```
def process_generic(rdd, mongo_url, rest_url):  
    count = rdd.count()  
    if count is 0:  
        return  
    count_packet_id = uuid.uuid4().hex  
    normalized_rdd = rdd.map(lambda e: repack(e, count_packet_id)).cache()  
    store_packets(count_packet_id, count, normalized_rdd, mongo_url)  
    signal_rest_server(count_packet_id,  
                        count,  
                        dict(normalized_rdd.map(  
                            lambda e: (e['service'], 1)).reduceByKey(add).collect()),  
                        rest_url)
```

Composing applications

- Dynamic resources
- Analytics components
- Custom glue
- Templating



Alice is building a data intensive _____ ,
noun
her data is stored in _____ , and it will be
IP address
_____ along the message queue named
verb
_____ .
proper noun

Bob maintains a _____ application to
adjective
display data, Alice _____ her data to Bob's
verb
server at _____ .
IP address

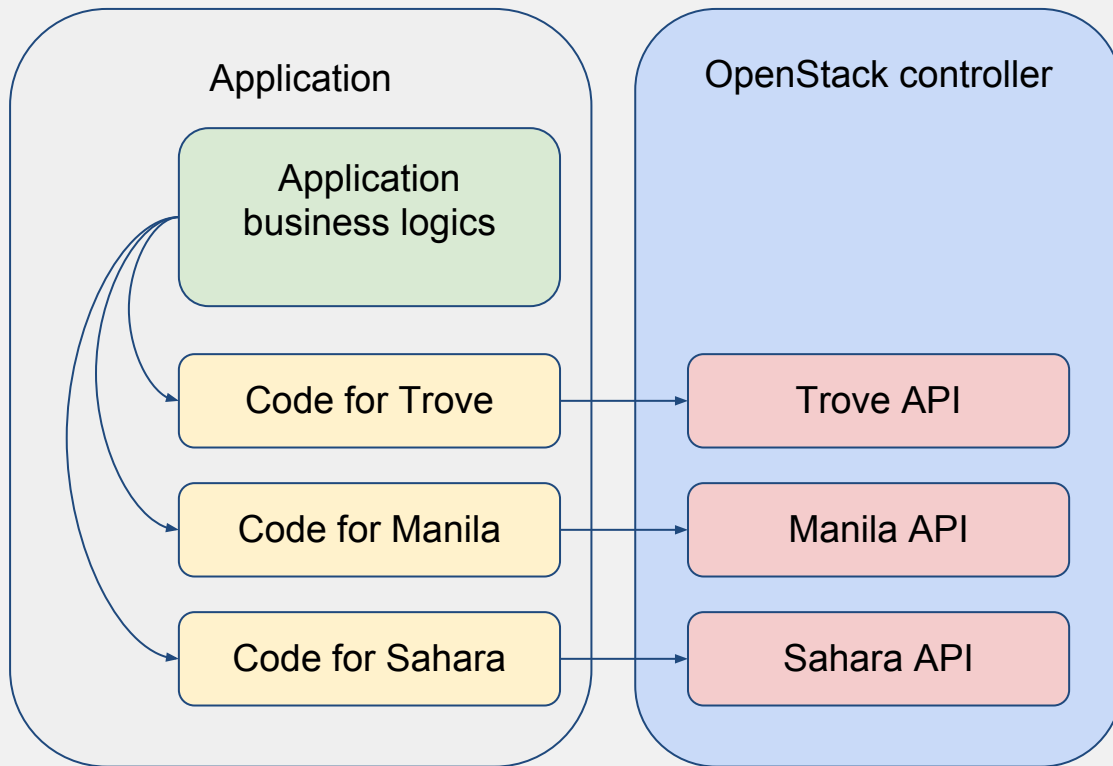
An alternative approach with Heat and Murano

- Provisioning all components through single endpoint
- Using generic syntax
- The syntax is declarative
- Higher level of API

Cloud client call workflow

- Application instantiates and authorizes OpenStack clients
- Application calls clients to reach OpenStack service on controller nodes
- Application handles clients' responses

Cloud client call workflow



Pros and cons

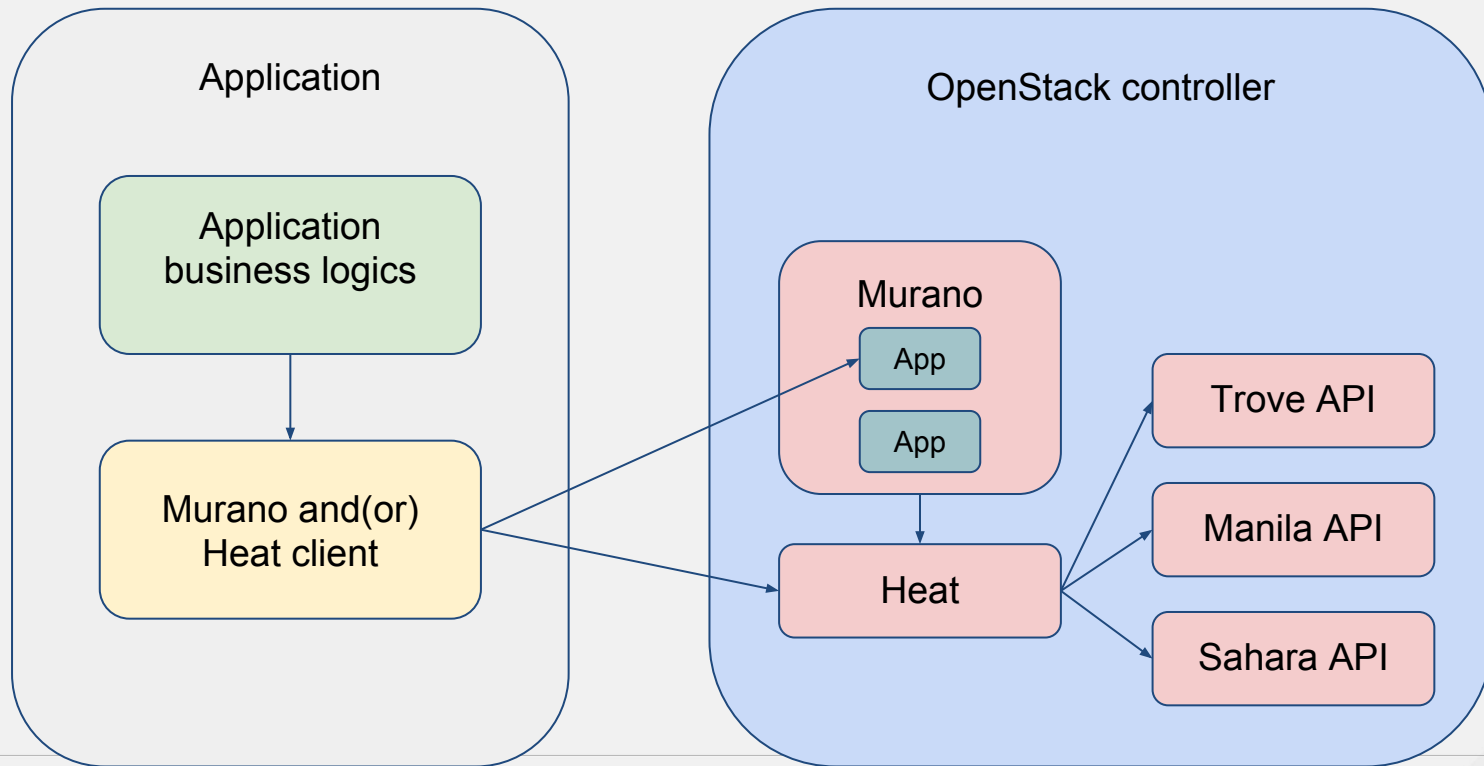
Advantages:

- Full control of the environment
- All operations may run completely independently

Disadvantages:

- Each client handle only a specific service
- You as a developer handle async waits, retries, errors, etc.

Heat and Murano approach



Pros and cons

Advantages:

- Access only one API endpoint
- Murano and heat catalogs of Apps
- UI integration into Horizon
- Handling of async ops, errors and retries is done by Heat and Murano

Disadvantages:

- Calls are limited to Murano apps and Heat templates

Building a Heat stack

- Declare parameters
- Declare resources. (using parameters)

Building a Heat template. Parameters.

Parameters are the customization point.

Input of basic types is allowed.

Example:

```
parameters:  
  image_id:  
    type: string  
    description: Image used for servers  
  instance_type:  
    type: string  
    default: m1.small  
  num_instances:  
    type: number  
    description: Number of instances to create  
    default: 1
```

Building a Heat template. Resources.

Resources section lists all VMs, Volumes, Shares, DBs, etc.

Parameters may be used here

Example:

resources:

server_group:

type: OS::Heat::InstanceGroup

properties:

LaunchConfigurationName: { get_resource: server_config }

AvailabilityZones: []

Size: { get_param: num_instances }

server_config:

type: AWS::AutoScaling::LaunchConfiguration

properties:

ImageId: { get_param: image_id }

InstanceType: { get_param: instance_type }

KeyName: { get_param: key_name }

Building a Murano App

- Create an App manifest
- Build an execution plan
- Describe deployment process
- Declare UI components
- Package everything

Building a Murano App. Manifest.

Format: 1.0

Type: Application

FullName: io.murano.apps.apache.ApacheHttpServer

Name: Apache HTTP Server

Description: |

The Apache HTTP Server Project is an effort to develop and maintain an open-source HTTP server for modern operating systems including UNIX and Windows NT.

...

Author: My Company, Inc

Tags: [HTTP, Server, WebServer, HTML, Apache]

Classes:

io.murano.apps.apache.ApacheHttpServer: ApacheHttpServer.yaml

Building a Murano App. Execution plan.

Name: Deploy Apache

Parameters:

enablePHP: \$enablePHP

Body: |

```
return apacheDeploy('{0}'.format(args.enablePHP)).stdout
```

Scripts:

apacheDeploy:

Type: Application

EntryPoint: runApacheDeploy.sh

Files: []

Options:

captureStdout: true

captureStderr: true

Building a Murano App. Deployment process.

Properties:

name:

Contract: \$.string().notNull()

enablePHP:

Contract: \$.bool()

instance:

Contract: \$.class(res:Instance).notNull()

Methods:

initialize:

Body:

<App initialization calls>

deploy:

Body:

<App deployment calls>

Building a Murano App. UI components


Forms:

- appConfiguration:
 - fields:
 - name: name
 - type: string
 - label: Application Name
 - initial: 'ApacheHttpServer'
 - description: >-
Enter a desired name for the application. Just A-Z, a-z, 0-9
 - name: enablePHP
 - label: Enable PHP
 - type: boolean
 - required: false

Packaging it all together

```
|_ Classes
|   |_ ApacheHttpServer.yaml
|
|_ Resources
|   |_ scripts
|       |_ runApacheDeploy.sh
|   |_ DeployApache.template
|
|_ UI
|   |_ ui.yaml
|
|_ logo.png
|
|_ manifest.yaml
```


Ready to go

demo ▾

Project ▾
Admin ▾
Identity ▾
Murano ^
Application Catalog ^
Environments
Applications
Manage ▾

Applications

Recent Activity




Apache HTTP Se...

The Apache HTTP Server Project is an effort to develop and maintain an open-source HTTP server


[Details »](#)

+ Add to Env

 Quick Deploy

The application life-cycle

- Iterating on projects
- Team-based building
- Development vs. production



Projects to watch

- OpenStack SDK
 - <http://developer.openstack.org/sdks/python/openstacksdk/>
- Murano
 - <https://murano.readthedocs.org/>
- Spark
 - <https://spark.apache.org>
- Big Data BBQ presentation examples
 - <https://github.com/elmiko/big-data-bbq>

Further reading

- OpenStack APIs and SDKs
 - <http://developer.openstack.org>
- Heat template Examples
 - <https://github.com/openstack/heat-templates>
- Application Catalog for Heat and Murano
 - <https://apps.openstack.org/>
- Guide for Murano Examples
 - https://murano.readthedocs.org/en/stable-liberty/draft/appdev-guide/step_by_step.html

Questions?



Get out there and build!